

Nishant Sharma

New York, NY · hellonishantsh@gmail.com · hellonish.dev · github.com/hellonish · linkedin.com/in/nishantsh20

Applied AI engineer and technical co-founder with 3 years of experience building production AI workflows, backend APIs, and customer-facing systems. Strong in Python, LangGraph, RAG, MCP-style tool integrations, FastAPI, Redis, PostgreSQL, AWS, and Docker. Shipped 12 client products and built Singularity, a multi-agent research platform deployed in 2 weeks.

Education

New York University, Tandon School of Engineering · New York, NY Master of Science in Computer Engineering · GPA: 3.7	May 2026
Dr. A.P.J. Abdul Kalam Technical University · Delhi, India Bachelor of Technology in Computer Science and Engineering · First Division with Distinction	Jun 2024

Technical Skills

Languages	Python · TypeScript · JavaScript · SQL · Java · C++
AI & Agents	LangGraph · RAG pipelines · multi-agent systems · LLM APIs · MCP · prompt engineering
Backend & APIs	FastAPI · Node.js · REST APIs · WebSockets · SSE · ARQ · Celery · Redis · async pipelines
Data & Storage	PostgreSQL · MySQL · MongoDB · Redis · AWS RDS · SQS · Qdrant · Chroma
Cloud/Infra	AWS EC2 · S3 · RDS · IAM · SQS · CloudFront · Docker · Kubernetes · Terraform · GitHub Actions · Caddy
Reliability	observability · structured logging · tracing · rate limiting · CI/CD · integration testing
Frontend	React · Next.js · TypeScript · Tailwind

Work Experience

Lead Software Engineer · Delhi, India · Ingelt Study Abroad	Sep 2023 – Aug 2024
- Led a 3-engineer backend team, owning task delegation, code reviews, PR approvals, sprint planning, and technical communication with company leadership. - Architected load-balanced AWS infrastructure with ALB, RDS, S3, and CloudFront, including async SQS pipelines and dead-letter handling for background jobs. - Built Jenkins CI/CD pipelines for automated test, build, and EC2 deployment, replacing manual releases with repeatable deployment workflows.	
Co-Founder and Chief Technology Officer · Delhi, India · Macverin Technologies · macverin.com	Jul 2022 – Aug 2023
- Sourced, scoped, built, and shipped 12 client platforms in 14 months, working directly with customers to translate business workflows into production software. - Built a visa consultancy operations platform covering inquiry intake, active case tracking, and meeting scheduling, replacing manual workflows for an 8-person team. - Designed a fare calculation layer between client pricing rules and Google Maps distance data, allowing operations teams to update pricing without changing integration code. - Led full-stack delivery across Node.js, FastAPI, React, Next.js, PostgreSQL, MySQL, MongoDB, and AWS, including deployment and post-launch support.	

Projects

Singularity - AI Deep Research Platform · singularity.hellonish.dev	Mar 2026 – Apr 2026
- Built and deployed a multi-agent research platform in 2 weeks, turning a manual research workflow into a single AI-assisted pipeline from query to cited report. - Designed a hierarchical agent system using LangGraph-style persistent state, DAG planning, and 44 modular skills across 11 domains for planning, retrieval, synthesis, and output generation. - Integrated MCP-style external tool connectivity and domain-specific skill routing, enabling agents to call retrieval, analysis, and output tools without hardcoded workflows - Built for unattended long-running jobs with Redis/ARQ workers, orphan job recovery, fail-open rate limiting, retries, and real-time SSE progress streaming. - Added structured logging, execution tracing, and integration tests against real PostgreSQL and Redis containers in GitHub Actions.	
Wand - AI Career Intelligence Platform · github	Jan 2026 – Apr 2026
- Built an AI workflow automation platform that reduces company research, JD analysis, contact strategy, and profile positioning from 30–60 minutes to one pipeline run. - Designed a 6-step concurrent agent pipeline using asyncio, Celery, Redis, and WebSockets, cutting wall-clock LLM latency roughly in half versus sequential execution. - Implemented live job-state streaming through Redis pub/sub so users can monitor long-running AI workflows step-by-step without polling. - Added BYOK, profile caching, structured output validation, and entity deduplication to reduce cost and make agent behavior measurable across runs.	
PGDR - Persistent Gradient Disagreement Reweighting · paper	Jan 2026 – May 2026
- Built reproducible PyTorch evaluation pipelines across ERM, JTT, and GroupDRO baselines using identical seeds, training budgets, and model selection criteria. - Improved Waterbirds worst-group accuracy from 68.38% to 81.52% while preserving 92.56% average accuracy across 3 seeds.	

Additional Experience

Machine Learning Teaching Assistant · NYU Tandon School of Engineering	Sep 2025 – May 2026
Led debugging sessions for 50+ graduate students on PyTorch training pipelines and standardized troubleshooting guides for recurring ML/environment failures.	