

Nishant Sharma

New York, NY · hellonishantsh@gmail.com · hellonish.dev · github.com/hellonish · linkedin.com/in/nishantsh20

Education

New York University, Tandon School of Engineering · New York, NY	May 2026
Master of Science in Computer Engineering · GPA: 3.7	
Dr. A.P.J. Abdul Kalam Technical University · Delhi, India	Jun 2024
Bachelor of Technology in Computer Science and Engineering · First Division with Distinction	

Technical Skills

Languages	Python · TypeScript · SQL · Java · C++
Frontend	React · Next.js · TypeScript · Tailwind CSS
Backend & APIs	FastAPI · Node.js · REST APIs · GraphQL · WebSockets · SSE · Celery · ARQ
AI & Agents	Multi-agent systems · RAG · LLM APIs · prompt engineering · tool calling · LangGraph
Data & Storage	PostgreSQL · Redis · MongoDB · MySQL · Qdrant · Chroma
Cloud & Infra	AWS · CloudFront · Docker · Kubernetes · Terraform · GitHub Actions
Observability/Testing	CI/CD · integration testing · structured logging · tracing · rate limiting · idempotency

Work Experience

Co-Founder and Chief Technology Officer · Delhi, India · Macverin Technologies · macverin.com	Jul 2022 – Aug 2023
- Delivered 12 client platforms in 14 months, working directly with customers on requirements, product scope, development, and launch. - Built customer-facing full-stack products across React, Next.js, TypeScript, Node.js, FastAPI, PostgreSQL, MongoDB, and AWS, supporting 1K+ daily users across 10+ clients. - Led client discovery, technical planning, deployment, and post-launch support, balancing product requirements, delivery speed, and production reliability.	
Lead Software Engineer · Delhi, India · Ingelt Study Abroad	Sep 2023 – Aug 2024
- Led a 3-engineer backend team, owning sprint planning, code reviews, PR approvals, and technical communication with company leadership. - Set up production AWS stack with load-balanced EC2, PostgreSQL on RDS, S3/CloudFront media delivery, and SQS-based background jobs for notifications and video uploads. - Built Jenkins CI/CD pipelines for automated testing, builds, and EC2 deployments, replacing manual releases with repeatable deployment workflows.	

Projects

Singularity - AI Deep Research Platform · singularity.hellonish.dev	Mar 2026 – Apr 2026
- Built and deployed a production multi-agent research platform in 13 days, shipping 25,500 LOC across a 6-container AWS stack with Next.js, FastAPI, PostgreSQL, Redis, and Docker. - Fine-tuned Qwen using RLVR for generating web search queries, raising internal eval accuracy of Qwen from 71% to 94% and outperforming GPT-5.1 (91%) on the same task set. - Designed agentic report generation with planning-first orchestration, targeted retrieval, source credibility checks, bottom-up synthesis, and final polish across 44 modular skills. - Engineered real-time AI workflow execution using Redis/ARQ workers, Redis pub/sub, FastAPI SSE streaming, retries, idempotency, and long-running job recovery.	
Wand - AI Career Intelligence Platform · github	Jan 2026 – Apr 2026
- Built an AI career intel platform that analyzes job descriptions, resumes, company context, and candidate profiles to produce match scores, contact strategy, and action plans. - Designed a 6-step concurrent LLM pipeline using asyncio, Celery, Redis, and WebSockets, cutting wall-clock latency roughly in half versus sequential execution. - Streamed live pipeline progress from backend workers to a React/TypeScript frontend using Redis pub/sub and WebSocket events, avoiding polling for long-running AI workflows. - Built a multi-source entity extraction and resolution pipeline that ingests profile data from resumes, LinkedIn exports, portfolios, and forms parses each source via LLM-structured extraction, deduplicates and priority-merges across sources, and emits a single validated canonical profile consumed by all downstream AI modules.	
PGDR - Persistent Gradient Disagreement Reweighting · paper	Jan 2026 – May 2026
- Proposed a label-free shortcut-bias detection method using last-layer gradient disagreement as a pseudo-group signal, improving Waterbirds worst-group accuracy from 68.38% to 81.52%. - Built reproducible PyTorch evaluation pipelines across ERM, JTT, GroupDRO, Group-Balanced, and Tail-loss baselines using shared seeds, training budgets, and model selection criteria. - Ran ablations diagnosing warmup quality and pseudo-group refresh stability as key failure modes for robustness.	